

**Title:** Energy Industry  
**Name:** Amandeep Singh  
**Student ID:** 921287533  
**Github Username:** Amandeep-Singh-24

| Checkpoint #   | Date Submitted |
|----------------|----------------|
| Checkpoint I   | 09/17/2023     |
| Checkpoint II  | 10/4/2023      |
| Checkpoint III | 10/25/2023     |

# Table of Contents

|  |           |
|--|-----------|
| <b>Cover Page.....</b>                             | <b>1</b>  |
| <b>Table of Contents.....</b>                      | <b>2</b>  |
| <b>Project Description.....</b>                    | <b>4</b>  |
| Database System for Energy Management (DBSEM)..... | 4         |
| Overview.....                                      | 4         |
| Unique Features.....                               | 4         |
| Use Cases.....                                     | 5         |
| Beneficiary Software Tools.....                    | 6         |
| <b>Functional Database Requirements.....</b>       | <b>8</b>  |
| User Management:.....                              | 8         |
| Energy Source Management:.....                     | 8         |
| Data Collection:.....                              | 8         |
| Analytics:.....                                    | 8         |
| Notifications and Alerts:.....                     | 9         |
| Renewable Energy Integration:.....                 | 9         |
| Billing and Payments:.....                         | 9         |
| Reporting:.....                                    | 9         |
| Third-party Integrations:.....                     | 9         |
| Feedback and Support:.....                         | 10        |
| Media Outreach:.....                               | 10        |
| Campaign:.....                                     | 10        |
| Partnership:.....                                  | 10        |
| Inventory:.....                                    | 11        |
| Supplier:.....                                     | 11        |
| Legal Case:.....                                   | 11        |
| Research Project:.....                             | 11        |
| Participant Feedback:.....                         | 11        |
| Training Program:.....                             | 11        |
| Employee:.....                                     | 12        |
| Asset:.....  | 12        |
| Facility:.....                                     | 12        |
| Aviation Analytics:.....                           | 12        |
| Power Metrics:.....                                | 12        |
| <b>Non-functional Database Requirements.....</b>   | <b>13</b> |
| 1. Performance.....                                | 13        |
| 2. Storage.....                                    | 13        |
| 3. Security.....                                   | 13        |
| 4. Scalability.....                                | 13        |
| 5. Reliability.....                                | 14        |

|  |           |
|--|-----------|
| 6. Interoperability.....                                     | 14        |
| 7. Maintainability.....                                      | 14        |
| <b>ERD Diagram.....</b>                                      | <b>15</b> |
| <b>Entity Description.....</b>                               | <b>16</b> |
| <b>Entity Establishment Relationship Diagram (EERD).....</b> | <b>19</b> |
| <b>Constraints Description.....</b>                          | <b>20</b> |

# Project Description

## Database System for Energy Management (DBSEM)

The continuous global shift towards sustainable energy is an increasing complexity of energy grids which prompts the necessity of an efficient, comprehensive, and adaptive database system progressing with the modern world. The Database System of Energy (DBSEM) aims to correct the inadequacies of the existing energy management databases, catering to the future of renewable energy sources, distributed grids, and adaptive energy pricing.

### Overview

DBSEM serves as a centralized repository that manages the information of various energy sources and integrates this data with smart grid technologies. This allows for energy storage solutions and an understanding of user consumption patterns, further optimizing energy usage. DBSEM serves as the bridge between energy providers and consumers or even regulators while sustaining a non-technical user interface that caters to a wide range of stakeholders.

### Unique Features

In today's rapidly evolving energy landscape, the need for agile and adaptive systems is more pronounced than ever before. Central to this transition is the introduction of the Database System for Energy Management (DBSEM). Traditional energy systems have long operated on static pricing models, leaving little room for the nuanced dynamics of real-time supply and demand fluctuations. With DBSEM, this paradigm is set to change. The system seamlessly incorporates **real-time data analysis**, paving the way for a genuinely **dynamic pricing model**. By assimilating variables such as instantaneous supply-demand curves, seasonal influences, and unforeseen tragic events, DBSEM ensures that pricing remains both adaptive and reflective of the current energy landscape. This dynamic approach is not just about economic efficiency; it's a pivotal step in promoting **sustainable energy practices**. With the ability to intricately monitor energy consumption, DBSEM can project precise household energy needs, ensuring conservation and optimal usage.

Furthermore, as we inch closer to a sustainable future, the importance of renewables cannot be understated. Recognizing this, DBSEM introduces **specialized modules tailored for renewable energy sources**. By leveraging the vast potential of solar and wind energy, the system integrates real-time data from weather forecasts, meticulously evaluating **energy storage capacities** against prevailing grid demands. This proactive approach ensures that renewables are not just supplemental but central to our energy strategies.

However, the ongoing decentralization of energy grids presents its own set of challenges. As the world increasingly gravitates towards **distributed and decentralized grid structures**, the need for an intelligent management system becomes imperative. Addressing this, DBSEM debuts a pioneering **topology mapping feature**. This tool empowers operators with unparalleled insights, enabling them to manage and optimize microgrids with a finesse previously unimagined. In essence, DBSEM isn't just another energy management tool; it's the bedrock on which the future of adaptive, sustainable, and decentralized energy solutions will be built.

---

## Use Cases

**Use Case:** Collaborative Energy Research & Development

**Actor:** R&D Scientist (Dr. Claire), Energy research institutions, DBSEM system

**Description:** Dr. Claire, leading a cutting-edge project on energy storage solutions, recognizes the importance of comprehensive data. The DBSEM offers her access to diverse energy patterns, new supplier technologies, and emerging renewable integration metrics. With this plethora of information, she's better equipped to draft research hypotheses and predict market needs. As her team iterates on prototypes, real-time updates from the DBSEM ensure the project remains relevant. Frequent collaboration with other research institutions, facilitated by shared DBSEM data, enhances their collective work quality. Their reliance on this system ensures that their solutions are both forward-thinking and deeply rooted in current industry challenges. Through this approach, Dr. Claire hopes to drive substantial positive shifts in the energy storage domain.

---

**Use Case:** Renewable Energy Certification

**Actor:** GreenCert Agency, Renewable energy producers, DBSEM system

**Description:** The GreenCert Agency ardently champions green energy practices. To maintain the integrity of their certification process, they've integrated with the DBSEM, pulling real-time data on energy production and carbon emissions. Energy producers vying for certification feed their metrics into the system, creating a transparent evaluation platform. The DBSEM's robust data analytics allow the GreenCert Agency to judge compliance against stringent global benchmarks. Producers that meet or exceed these standards are awarded the GreenCert, distinguishing them in a competitive marketplace. By maintaining this rigorous and transparent process, the GreenCert Agency hopes to incentivize more producers to adopt sustainable practices, shaping a greener energy future.

---

**Use Case:** Industrial Energy Optimization

**Actor:** Factory Manager (Raj), Machinery vendors, DBSEM system

**Description:** For Raj, every ounce of energy saved in his factory translates to enhanced operational efficiency. Integrating his factory setup with the DBSEM, he's empowered with real-time insights into the energy consumption patterns of each production line. Drawing

correlations between energy data and production output, he identifies areas ripe for optimization. Regular consultations with machinery vendors allow for rapid interventions and retrofitting. As Raj introduces new machinery or modifies production strategies, continuous feedback from the DBSEM ensures he's on the right energy-saving track. His commitment to the system not only ensures the factory's sustainability but also significantly boosts its bottom line, making Raj a pioneering figure in his industrial circle.

---

**Use Case:** City-wide Energy Conservation Initiatives

**Actor:** City Planner (Linda), Residents, DBSEM system

**Description:** Linda's ambitious vision for her city is grounded in sustainability. With the DBSEM, she harnesses data from various city sectors, piecing together a holistic view of energy consumption. Identifying energy-intensive zones, she meticulously designs targeted conservation initiatives. By sharing these insights with city residents through interactive DBSEM dashboards, she instills a collective energy consciousness. As weeks turn into months, Linda observes positive shifts in city-wide energy metrics. Through regular DBSEM-generated reports, she validates and celebrates these successes, keeping the momentum alive. Her endeavors inspire other cities, showcasing the transformational power of data-driven conservation initiatives.

---

**Use Case:** Electric Vehicle Charging Network

**Actor:** EV Driver (Pedro), Charging station operators, DBSEM system

**Description:** Pedro's electric vehicle (EV) isn't just a mode of transport—it's a statement of his environmental commitment. The DBSEM becomes his co-pilot, offering real-time updates on charging station availability, dynamic energy rates, and peak charging times. This information ensures that his long drives are punctuated with efficient charging stops, maximizing his EV's performance. On the flip side, charging station operators harness the DBSEM's insights to adjust energy supply based on demand. By dynamically pricing energy during off-peak times, they encourage balanced charging across their network. This synergy between Pedro and the operators, facilitated by the DBSEM, ensures a seamless EV charging experience. As more drivers like Pedro take to the roads, the DBSEM's role becomes ever more critical in shaping an efficient EV charging ecosystem.

---

## **Beneficiary Software Tools**

In the realm of building energy and lighting system simulations, **EnergyPlus** stands out as a foremost software solution. It boasts capabilities that go beyond traditional modeling, making it invaluable to modern architectural undertakings. However, its potential multiplies when combined with DBSEM's sophisticated energy management features. By utilizing DBSEM's **dynamic pricing models** and comprehensive **grid analytics**, EnergyPlus can venture into unprecedented realms of realism. This integration equips architects and engineers with the tools to simulate future energy costs with remarkable precision, allowing for better-informed decisions in their projects. On the other hand, **OpenDDS**, a renowned power flow application, finds its

strengths amplified with DBSEM's insights. The intricate **microgrid management** and **renewable integration features** of DBSEM seamlessly complement OpenDDS's framework. Such collaboration ensures that OpenDDS is not only equipped to simulate the present-day grids but is also future-ready, embodying the evolving essence of modern energy landscapes.

# Functional Database Requirements

## **User Management:**

- 1.1. Users shall be able to register with a unique email and password.
- 1.2. Users shall be able to log in using their email and password.
- 1.3. Users can recover their password through email authentication.
- 1.4. Users can edit and update their profile information.
- 1.5. Users can delete their accounts.

## **Energy Source Management:**

- 2.1. Users shall be able to add multiple energy sources to their profile.
- 2.2. Each energy source shall have a unique identification.
- 2.3. Users can edit and update the details of their energy sources.
- 2.4. Users can delete an energy source from their profile.
- 2.5. Users can share their energy source data with other users.
- 2.6. Users can view analytics of each energy source separately.
- 2.7. The system should allow categorization of energy sources (e.g., solar, wind, hydro).
- 2.8. Users can set primary and secondary energy sources for prioritized consumption.

## **Data Collection:**

- 3.1. An energy source shall record its consumption every hour.
- 3.2. Data collection should support real-time integration with smart meters.
- 3.3. Manual data entry should be available for non-integrated energy sources.
- 3.4. Users shall be able to view historical consumption data for each energy source.
- 3.5. The system shall support integration with IoT devices for more accurate data collection.
- 3.6. Users can set the frequency of data collection other than hourly, such as every 30 minutes.

## **Analytics:**

- 4.1. Users shall be provided with daily consumption analytics.
- 4.2. Weekly and monthly consumption patterns should be available.
- 4.3. The system shall provide peak consumption times for each day, week, and month.
- 4.4. The system shall provide predictions for the next day's consumption.
- 4.5. Users shall receive recommendations for energy conservation based on their consumption patterns.
- 4.6. Users can compare their energy consumption with the average consumption of similar users.
- 4.7. The system should predict the best time for energy consumption based on the user's behavior.
- 4.8. Users can view graphical representations of their energy source outputs.



**Notifications and Alerts:**

- 5.1. Users shall receive alerts on consumption target breaches.
- 5.2. Peak demand prediction alerts shall be provided.
- 5.3. Users shall receive maintenance notifications for their energy equipment.
- 5.4. Carbon footprint tracking alerts shall be sent monthly.
- 5.5. Users can customize the type and frequency of notifications they receive.
- 5.6. Alerts for unusual or unexpected energy consumption spikes shall be provided.

**Renewable Energy Integration:**

- 6.1. Solar and wind energy production data can be integrated.
- 6.2. Users shall be able to view their renewable energy contribution.
- 6.3. The system shall provide optimal times for renewable energy consumption.
- 6.4. Users shall receive recommendations for optimizing energy storage based on renewable sources.
- 6.5. Users can integrate energy storage systems like batteries.
- 6.6. The system will provide alerts when renewable sources are producing energy in excess of storage capabilities.

**Billing and Payments:**

- 7.1. Users can view their monthly energy bills.
- 7.2. Users can make payments for their energy bills through the system.
- 7.3. A history of all payments shall be maintained.
- 7.4. The system shall provide predictions of future bills based on current consumption patterns.
- 7.5. Users can set up automatic bill payments.
- 7.6. The system should support integration with multiple payment gateways.
- 7.7. Users can view a detailed breakdown of their energy costs.

**Reporting:**

- 8.1. Users can generate and download energy consumption reports.
- 8.2. The system shall provide customizable report templates.
- 8.3. Reports can be generated daily, weekly, monthly, or for custom periods.
- 8.4. The system shall provide options for visual representation (graphs, pie charts) in the reports.
- 8.5. Users can schedule automatic report generation at specified intervals.
- 8.6. Reports can include user-specific notes or annotations.

**Third-party Integrations:**

- 9.1. The system should allow for integration with home automation systems.
- 9.2. Integration with SCADA systems for industrial users should be possible.
- 9.3. APIs should be provided for third-party developers to create custom integrations.

- 9.4. The system should support integration with energy trading platforms for users who want to sell excess energy.
- 9.5. The system should be compatible with major operating systems and devices.

#### **Feedback and Support:**

- 10.1. Users shall be able to provide feedback on the system.
- 10.2. A support ticketing system shall be in place for user queries.
- 10.3. Users can view the status of their support tickets.
- 10.4. The system shall provide a knowledge base for frequently asked questions.
- 10.5. The system shall provide options for live chat support.
- 10.6. Support should also be available via phone for critical issues.

#### **Mailing:**

- 11.1 Mailing shall have a unique MailID.
- 11.2 Mailing shall have an address associated with it.
- 11.3 Mailing may have one or more delivery dates.
- 11.4 Mailing can be associated with zero, one, or many users through the mailing list.

#### **Media Outreach:**

- 12.1. A media outreach shall have a unique identifier within the context of a specific campaign.
- 12.2. A media outreach shall always be associated with a campaign.
- 12.3. A media outreach shall have a content type associated with it.
- 12.4. A media outreach shall specify a target audience.
- 12.5. A media outreach shall specify a platform on which the outreach is executed.
- 12.6. The combination of content type, target audience, and platform may be used to differentiate between different media outreaches within a campaign.

#### **Campaign:**

- 13.1. A campaign shall have a unique strong campaign ID.
- 13.2. A campaign shall have a campaign name associated with it.
- 13.3. A campaign shall specify a target audience.
- 13.4. A campaign shall have a composite date attribute, which may consist of attributes like start date, end date, etc.
- 13.5. A campaign can be associated with zero, one, or many users.

#### **Partnership:**

- 14.1. A partnership shall have a unique partnershipID.
- 14.2. A partnership shall specify a duration.
- 14.3. A partnership shall have a start date.
- 14.4. A partnership shall specify a partner name.

14.5. A partnership shall specify a partner type.

**Inventory:**

15.1. An inventory record supplied by a specific supplier shall have a unique inventoryID.

15.2. An inventory record shall specify a quantity.

15.3. An inventory record shall specify a product name.

15.4. An inventory record must be associated with a specific supplier.

**Supplier:**

16.1. A supplier shall have a unique supplier ID.

16.2. A supplier shall specify a supplier name.

16.3. A supplier shall specify a product type.

**Legal Case:**

17.1. A legal case shall have a unique caseid.

17.2. A legal case shall specify a status.

17.3. A legal case shall specify the lawyer assigned to it.

**Research Project:**

18.1. A research project shall have a unique projected identifier.

18.2. A research project shall specify a project name.

18.3. A research project shall specify a funding source.

18.4. A research project shall specify start and end dates.

**Participant Feedback:**

19.1. Participant feedback shall have a unique feedback sequence within its associated context (e.g., a specific event).

19.2. Participant feedback shall capture comments.

19.3. Participant feedback shall specify the fullname of the participant providing the feedback.

**Training Program:**

20.1. A training program shall have a unique training id.

20.2. A training program shall specify its duration.

20.3. A training program shall specify a training name.

**Employee:**

- 21.1. An employee shall have a unique employeeid.
- 21.2. An employee shall specify a fullname.
- 21.3. An employee shall specify a position.

**Asset:**

- 22.1. An asset shall have a unique assetid.
- 22.2. An asset shall specify its purchase date.
- 22.3. An asset shall specify its value.
- 22.4. An asset shall specify its type.

**Facility:**

- 23.1. A facility shall have a unique facilityID.
- 23.2. A facility shall specify its usage type.
- 23.3. A facility shall specify its size.
- 23.4. A facility shall specify its location.

**Aviation Analytics:**

- 24.1. An aviation analytic record shall specify the fuel cost.
- 24.2. An aviation analytic record shall specify the time of the data.
- 24.3. An aviation analytic record shall specify the aircraft type.

**Power Metrics:**

- 24.1. A power metric record shall specify infrastructure plans.
- 24.2. A power metric record shall specify the average cost.
- 24.3. A power metric record shall specify the total power generated.

# Non-functional Database Requirements

## 1. Performance

- 1.1 The database system shall support concurrent transactions, allowing multiple users to execute tasks without conflicts.
- 1.2 The response time for standard queries shall not exceed 2 seconds under normal load conditions.
- 1.3 The database should be able to handle a minimum of 10,000 simultaneous user requests without performance degradation.

## 2. Storage

- 2.1 The database system shall assign 10 MB of memory per table for optimized data retrieval.
- 2.2 The database system should support persistent storage, ensuring that data remains intact even after system shutdowns or restarts.
- 2.3 The system shall automatically compress older data to ensure efficient utilization of storage space.

## 3. Security

- 3.1 Only encrypted passwords shall be supported by the database system, following the SHA-256 encryption standard or better.
- 3.2 All values inserted into the database shall be consistent with the attribute's datatype and domain to prevent data inconsistency.
- 3.3 The database shall be automatically backed up every day at 11:59 p.m., with backups stored in a geographically distant location for redundancy.
- 3.4 Role-based access control (RBAC) shall be implemented, ensuring users have access only to the data and functions they are authorized for.

## 4. Scalability

- 4.1 The database shall be designed for both vertical and horizontal scalability to cater to future growth and expansion.
- 4.2 The system should support the addition of new tables, views, or other structures without affecting existing functionality.

## **5. Reliability**

- 5.1 The database shall offer 99.99% uptime, excluding planned maintenance windows.
- 5.2 Failover mechanisms shall be implemented to minimize system downtime during unexpected outages.

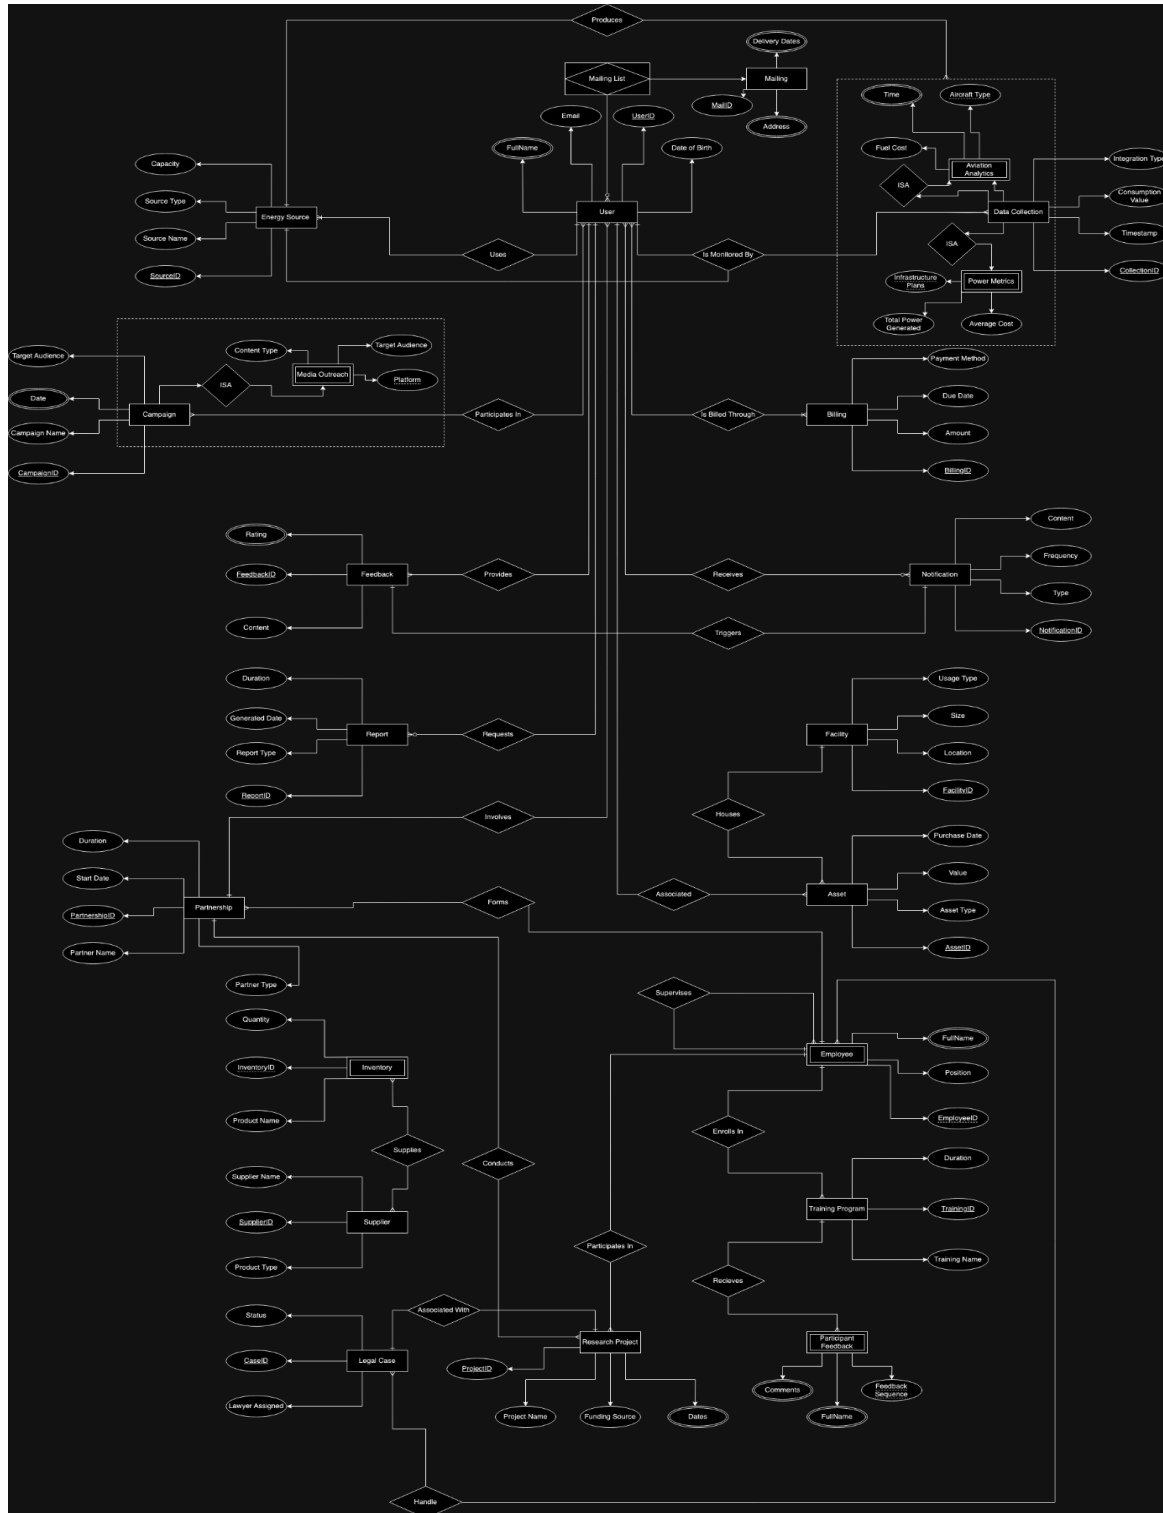
## **6. Interoperability**

- 6.1 The database system shall provide APIs for integration with third-party applications and platforms, including energy management tools and analytics software.
- 6.2 It shall support multiple database drivers and connectivity options, such as JDBC and ODBC, for diverse application support.

## **7. Maintainability**

- 7.1 The system shall support easy patching and updates without significant downtime.
- 7.2 Database logs and error reports shall be generated automatically, aiding in diagnostics and troubleshooting.

# ERD Diagram



# Entity Description

## 1. User (Strong):

- UserID: key, numeric
- Email: simple, alphanumeric
- FullName: multivalue, alphanumeric
- Date of Birth: simple, numeric

## 2. Energy Source (Strong):

- Capacity: simple, numeric
- Source Type: simple, alphanumeric
- Source Name: simple, alphanumeric
- SourceID: key, numeric

## 3. Campaign (Strong):

- Date: composite, timestamp
- Target Audience: simple, alphanumeric
- Campaign Name: simple, alphanumeric
- Campaign ID: key, numeric

## 4. Media Outreach (Weak):

- Content Type: simple, alphanumeric
- Target Audience: simple, alphanumeric
- Platform: key, numeric

## 5. Feedback (Strong):

- Rating: composite, alphanumeric
- FeedbackID: key, numeric
- Content: simple, alphanumeric

## 6. Report (Strong):

- Duration: simple, numeric
- Generated Date: simple, numeric
- Report Type: simple, alphanumeric
- ReportID: key, numeric

## 7. Partnership (Strong):

- Duration: simple, numeric
- Start Date: simple, numeric
- PartnershipID: key, numeric
- Partner Name: simple, alphanumeric



**8. Inventory (Weak):**

- Quantity: simple, numeric
- IneventoryID: key, numeric
- Product Name: simple, alphanumeric

**9. Supplier (Strong):**

- Supplier Name: simple, alphanumeric
- SupplierID: key, numeric
- Product Type: simple, alphanumeric

**10. Legal Case (Strong):**

- Status: simple, alphanumeric
- CaseID: key, numeric
- Lawyer Assigned: simple, alphanumeric

**11. Research Project (Strong):**

- ProjectID: key, numeric
- Project Name: simple, alphanumeric
- Funding Source: simple, alphanumeric
- Dates: composite, timestamp

**12. Participant Feedback (Weak):**

- Comments: composite, alphanumeric
- Fullname: composite, alphanumeric
- Feedback Sequence: key, alphanumeric

**13. Training Program (Strong):**

- Duration: simple, alphanumeric
- TrainingID: key, numeric
- Training Name: simple, alphanumeric

**14. Employee (Weak):**

- FullName: composite, alphanumeric
- Position: simple alphanumeric
- EmployeeID: key, numeric

**15. Asset (Strong):**

- Purchase Date: simple, alphanumeric
- Value: simple, alphanumeric
- Asset Type: simple, alphanumeric
- AssestID: key, numeric

**16. Facility (Strong):**

- Usage Type: simple, alphanumeric
- Size: simple, numeric
- Location: simple, alphanumeric
- FacilityID: key, numeric

**17. Notification (Strong):**

- Content: simple, alphanumeric
- Frequency: simple, numeric
- Type: simple, alphanumeric
- NotificationID: key, numeric

**18. Billing (Strong):**

- Payment Method: simple, alphanumeric
- Due Date: simple, timestamp
- Amount: simple, numeric
- BillingID: key, numeric

**19. Data Collection (Strong):**

- Integration Type: simple, alphanumeric
- Consumption Value: simple, numeric
- Timestamp: simple, numeric
- CollectionID: key, numeric

**20. Aviation Analytics (Weak):**

- Aircraft Type: key, alphanumeric
- Fuel Cost: simple, numeric
- Time: composite, numeric

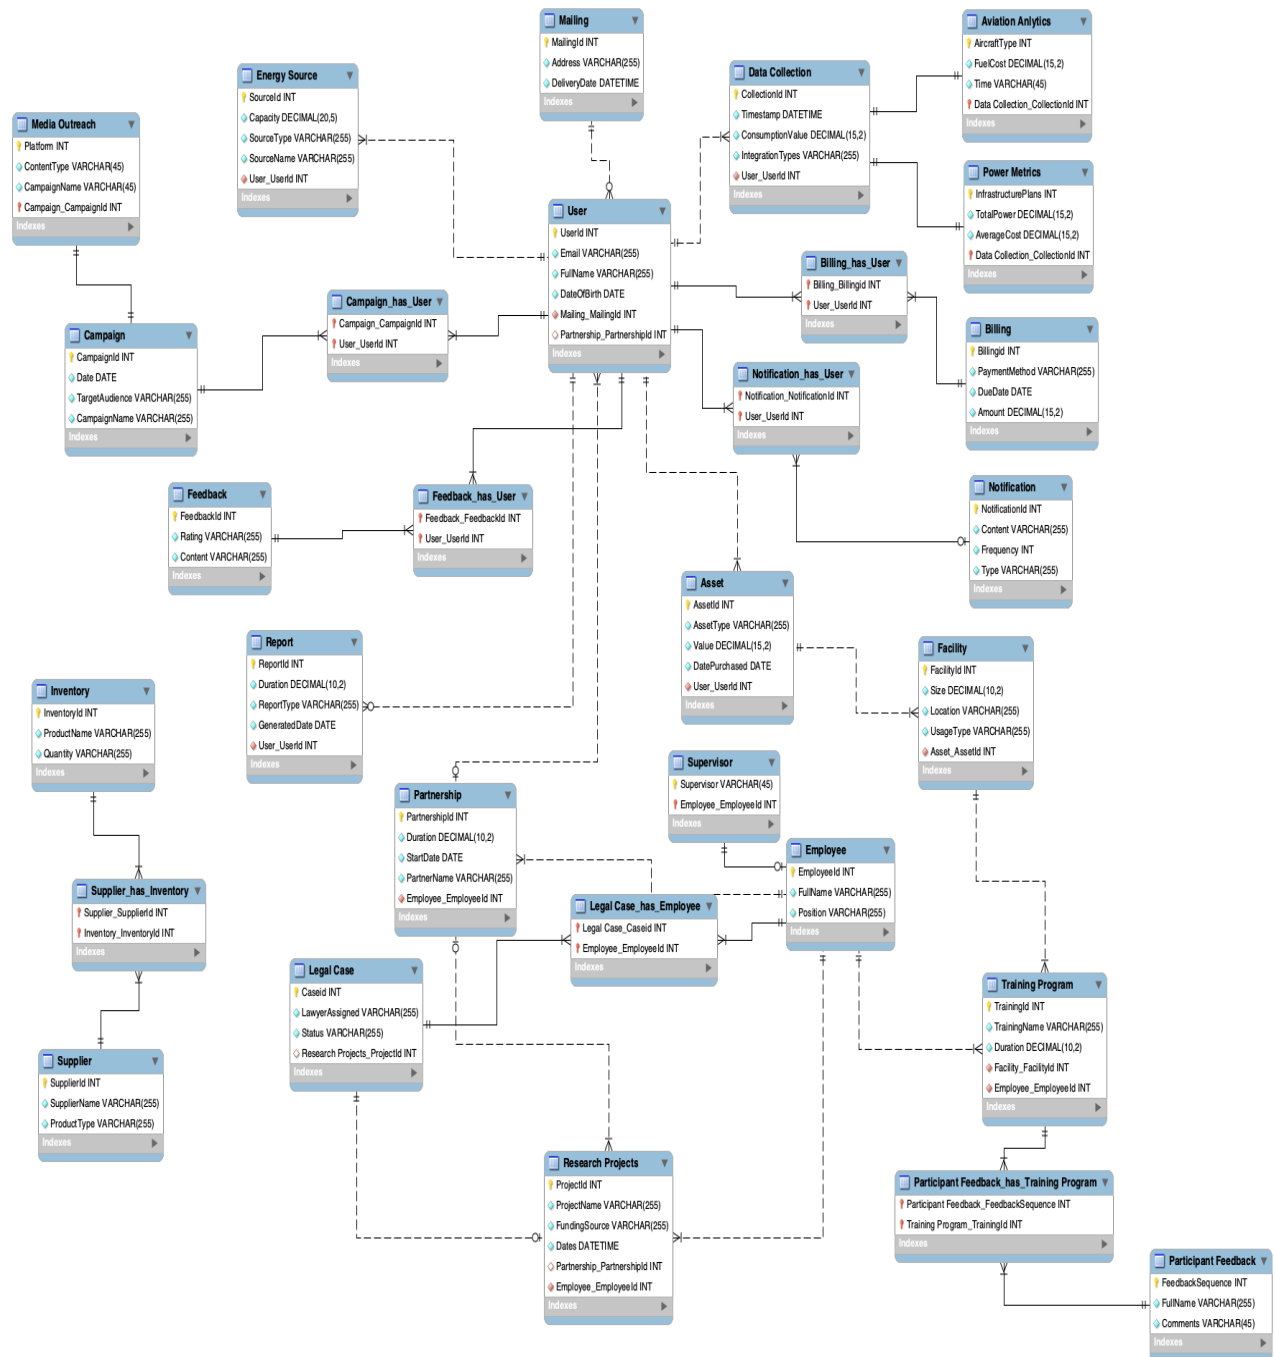
**21. Power Metrics (Weak):**

- Infrastructure Plans: key, alphanumeric
- Total Power Generated: simple, numeric
- Average Cost: simple, numeric

**22. Mailing (Strong):**

- Delivery Dates: composite, timestamp
- MailID: key, numeric
- Address: composite, alphanumeric

# Entity Establishment Relationship Diagram (EERD)



## Constraints Description

| Table      | FK         | ON DELETE | ON UPDATE | Comments  |
|------------|------------|-----------|-----------|---|
| Asset      | UserId     | SET NULL  | CASCADE   | Assets don't disappear when users leave. But if a user's info is updated, it should reflect on assets they're related to.                           |
| User       | MailingId  | SET NULL  | CASCADE   | Users might opt out of mailings, but we don't want to delete the user. Updates to user info should reflect on their mailing settings.               |
| Supervisor | EmployeeId | SET NULL  | CASCADE   | Supervisors might change roles, but employees under them don't get removed. Updates to supervisor info should reflect on associated employees.      |
| Employee   | TrainingId | SET NULL  | CASCADE   | If an employee leaves, we don't want to lose their training records. However, updates to employee details should reflect on their training records. |
| Facility   | AssetId    | SET NULL  | CASCADE   | Facilities can change, but assets inside  |

|                 |            |          |         |  |
|-----------------|------------|----------|---------|--|
|                 |            |          |         | them persist. Changes in facility details should cascade to assets located there.  |
| TrainingProgram | FeedbackId | CASCADE  | CASCADE | Feedback is tightly tied to a training program. If a training program is deleted, related feedback should be removed. Similarly, updates should reflect immediately. |
| Inventory       | SupplierId | CASCADE  | CASCADE | The link between inventory and supplier is tight. If an inventory item is removed, supplier relations should also be updated or removed.                             |
| LegalCase       | EmployeeId | CASCADE  | CASCADE | A legal case might be tied to the performance or involvement of an employee. Thus, updates or deletions should be mirrored.  |
| ResearchProject | Caseld     | SET NULL | CASCADE | While research influences cases, a project's conclusion shouldn't delete associated cases. But updates should  |

|                |            |          |         |  |
|----------------|------------|----------|---------|--|
|                |            |          |         | cascade.   |
| Partnership    | UserId     | SET NULL | CASCADE | Users might leave a partnership, but the partnership remains. However, updates to user info should reflect on their partnership relations. |
| Billing        | UserId     | CASCADE  | CASCADE | If a user is removed, their billing data should also be deleted. Updates in user data should instantly reflect in billing.                 |
| DataCollection | AviationId | SET NULL | CASCADE | Data might go obsolete, but aviation records persist. Yet, updates in data should reflect on associated aviation records.                  |
| Notification   | UserId     | CASCADE  | CASCADE | Notifications are user-specific. If a user is removed or their details change, it should be reflected in their notifications.              |